Philosopher's Stone: A Scope-Relative, Diachronic System for Modeling Philosophical Maturity and Epistemic Stability

Raul Miguel P. Paras III

April 2026

Abstract

Philosopher's Stone is a browser-native system for modeling worldview structure, philosophical maturity, and epistemic stability by separating probabilistic extraction from deterministic scoring and geometric projection. The system is built around the Epistemic Octahedron, a surface-based model whose three semantic axes represent empathy versus practicality, wisdom versus knowledge, and negative versus positive epistemic stability. This paper presents the current system as it now exists conceptually and operationally: a constrained extraction contract, a deterministic profiler, a scope-sensitive maturity model, a diachronic profile state, and a geometric visualizer that preserves the octahedral surface rule. The central claim of the paper is that objective peak philosophical maturity should be reachable whenever a text fully covers the scope it opens and genuinely integrates all four lateral dimensions with positive epistemic stability. Peak is therefore not treated as a lifetime prestige marker reserved for broad biography or old age; it is a scope-relative completion state. Later texts may widen scope and lower the profile without making the earlier completion false. This paper also outlines the system's scientific rationale, implementation logic, validation path, limitations, and design choices needed to make the system public, inspectable, and falsifiable rather than mystical, purely intuitive, or black-boxed.

1 Introduction

Public disagreement is often shallow for two different reasons at once. First, people frequently argue over conclusions without a shared model of what counts as genuine philosophical maturity. Second, many current AI-assisted systems either collapse into vague moral language or hide their logic behind opaque scoring and inaccessible internal state. Philosopher's Stone is an attempt to solve both problems by making the whole stack inspectable.

The system does not ask an LLM to "decide who is wise." It asks an LLM to extract structured evidence under a fixed contract. A deterministic profiler then aggregates that evidence into semantic values, computes maturity and stability, and projects the resulting worldview to a surface-constrained geometric model. The visualizer only renders the output of that deterministic layer. This separation is the core architectural discipline of the project.

The present paper is a system paper rather than a graph-theory paper. The graph theory belongs to the Epistemic Octahedron itself. The purpose here is different: to explain how the current pipeline

works, what problem it is trying to solve, what changes were needed to align it with the theory, and how the system can be validated scientifically rather than defended rhetorically.

2 Theoretical Basis and System Goal

The Epistemic Octahedron models active worldviews on the surface

$$|x| + |y| + |z| = 1. \tag{1}$$

The semantic axes are:

$x < 0$	practicality,	$x > 0$	empathy,
$z < 0$	knowledge,	$z > 0$	wisdom,
$y < 0$	negative epistemic stability,	$y > 0$	positive epistemic stability.

The crucial conceptual distinction is between the null origin, the lower vertex, and the upper vertex. The origin is pre-philosophical nullity: no active worldview has formed strongly enough to be plotted on the surface. The lower vertex is epistemic collapse: an active worldview that is maximally negatively stable. The upper vertex is objective peak philosophical maturity: the state in which empathy, practicality, wisdom, and knowledge have all been genuinely encountered and integrated without passive destabilization by asymmetry.

That last clause matters. The upper vertex is not defined by flat sameness or the absence of lateral content. It is defined by considered integration. This means a mature philosophy can still show emphasis. A text can lean practical or wisdom-heavy while remaining mature if the opposite poles are still meaningfully considered. The system therefore needs to distinguish emphasis from neglect and scope-complete maturity from unresolved incompleteness.

The goal of Philosopher's Stone is to operationalize that distinction in a way that is public, inspectable, and capable of diachronic use.

3 System Overview

The current system has five layers:

1. user text input,
2. structured extraction under a fixed JSON contract,
3. deterministic profiling and semantic aggregation,
4. projection to the octahedron surface,
5. interactive visualization and profile inspection.

The LLM is not allowed to compute the final plot. It extracts support, confidence, local signals, gate events, scope information, dimension consideration, and profile-update signals. The profiler then computes the final semantic parameters and projected coordinates. This design is not cosmetic. It is the mechanism that keeps the system auditable.

3.1 Why the separation matters

If the LLM were allowed to output final coordinates directly, the system would become difficult to benchmark, difficult to validate, and difficult to debug. Any error in the final plot would be ambiguous: was it a theory problem, a prompt problem, a scoring problem, or a hallucination? By splitting extraction from scoring, the system can inspect intermediate structure and localize failure.

The separation also allows model benchmarking. Two LLMs can be compared on extraction quality while holding the profiler fixed. Likewise, profiler rules can be stress-tested while holding the extractor contract fixed.

4 Current Extraction Contract

The current contract is an evidence-extraction schema, not a verdict engine. It requires the model to return, among other fields:

- a semantic grid with support and confidence for empathy, practicality, wisdom, knowledge, x-integration, z-integration, positive epistemic stability, and negative epistemic stability,
- dimension consideration for all four lateral dimensions,
- claim commitments and scope effects,
- a scope profile including claimed scope, scope completion, scope expansion, unresolved gaps, and relevant versus irrelevant gates,
- local extraction fields such as principles, boundaries, claimed values, tradeoffs, and contradictions,
- axis events,
- local positive and negative epistemic signals,
- triggered gate events and gate-update proposals,
- profile-update signals and canon-optimization suggestions.

Several design choices are especially important.

4.1 Target attribution

The contract distinguishes the profiled target from criticized systems, described others, quoted views, and cautionary examples. This prevents the system from confusing "the author criticizes an irrational institution" with "the author is irrational." It also prevents the graph from treating outside-target failures as self-failures.

4.2 Dimension consideration is not crude sentiment

A dimension can be directly engaged, acknowledged, tradeoff-engaged, explicitly deprioritized, explicitly rejected, or not evidenced. This is important because criticizing a distorted use of empathy is not identical to rejecting empathy as a dimension. The system must preserve that distinction if it is to model philosophical maturity rather than surface tone.

4.3 Scope is explicit

The system now requires a `scope_profile` every time. This makes the model state what scope is being claimed, whether the text covers the territory it opens, whether the current text widens scope, and which gates are relevant or irrelevant within that scope. This is the structural basis for scope-relative maturity.

4.4 Contract role

The contract should be lean where possible, but not so lean that semantic distinctions collapse. Type and formatting expectations can often be carried by schema names; interpretive invariants still need explicit rules. This keeps the extractor constrained without turning the contract into a bloated instruction dump.

5 Deterministic Profiler

The profiler is the mathematical authority of the system. It normalizes the extracted structure, merges it into profile state, updates gate states, computes semantic parameters, and projects the result to the octahedral surface.

5.1 Semantic parameters

The system computes semantic parameters (a, b, s) before final projection. Here a represents the empathy–practicality axis, b represents the wisdom–knowledge axis, and s represents epistemic

stability. A representative form of the current calculation is:

$$a = \mathrm{clamp}\left(\frac{(E-P)(1-\rho_x\lambda_I)}{\sigma_x}, -1, 1\right), \tag{2}$$

$$b = \mathrm{clamp}\left(\frac{(W-K)(1-\rho_z\lambda_I)}{\sigma_z}, -1, 1\right), \tag{3}$$

$$s = \mathrm{clamp}\left(\frac{Y^+ - Y^- + B_I - P_U - \alpha C}{\sigma_y} + \beta G^+ - \gamma G^-, -1, 1\right). \tag{4}$$

The exact internal ingredients evolve with implementation, but the structure is stable: lateral semantic mass, integration bonuses, contradiction penalties, and gate influences are aggregated deterministically rather than guessed by the extractor.

The active worldview is then projected to the surface by

$$(x, y, z) = \frac{(a, s, b)}{|a| + |s| + |b|} \quad \text{when} \quad |a| + |s| + |b| > 0. \tag{5}$$

If the semantic mass is truly null, the system returns the origin.

5.2 Gate structure

The profiler maintains six meta-epistemic gates:

G1 counter-consideration,

G2 non-strawman,

G3 self-correction,

G4 contradiction handling,

G5 reality contact,

G6 non-self-sealing.

These are not decorative labels. They are part of the epistemic stability calculation and are stored diachronically in gate state. The current system tracks not only gate scores but also positive and negative event counts, last evidence spans, and persistent status markers.

5.3 Current profile state

The profiler also maintains a cumulative profile state consisting of core principles, core boundaries, meta-epistemic markers, risk notes, and a gate snapshot. New entries can refine or reinforce the profile rather than replacing it wholesale.

6 Scope-Relative Peak Maturity

A major correction to the system was moving away from a maturity model that effectively pushed almost every text away from the peak. That older behavior treated peak like an asymptote or an exclusive lifetime badge. The current design is different.

6.1 Peak is a completion state

Objective peak philosophical maturity is now treated as a reachable completion state relative to scope. If a text opens a scope and genuinely covers that scope while integrating empathy, practicality, wisdom, and knowledge under positive epistemic stability, the system should allow that text to reach the peak.

This means peak is not reserved for age, biography, or sheer quantity of life topics. A text on a single domain can hit peak if that domain is fully and maturely handled. For example, a philosophy of dating advice could be peak-worthy within that domain if it truly integrates all four dimensions and remains epistemically sound. Peak is therefore not exclusive to broad adult life review.

6.2 Why this matters

Without scope-relative peak, the graph sends the wrong message. It encourages a view of maturity as something permanently withheld. That contradicts the theory. A mature scoped philosophy should be allowed to show itself as mature.

6.3 Diachronic expansion

The system is designed to be diachronic. A profile can later widen its active scope. If a new text introduces unresolved territory, falsehood, contradiction, or gate failure, the graph can fall from the peak. This does not mean the earlier peak was fake. It means the active profile now covers a larger domain, and that enlarged domain is not yet equally mature.

This is one of the most important design features of the whole system. It allows the graph to function both as a representation of present maturity and as an instrument for self-testing, challenge, revision, and group learning.

7 Why the Current Design Fits the Theory Better

The system now aligns much more closely with the theory in four ways.

7.1 First, it distinguishes emphasis from neglect

A philosophy can lean practical, empathetic, wise, or knowledge-heavy without being immature. Immaturity appears when the neglected poles are not really considered or when the system becomes

epistemically unstable, not merely when one side is emphasized.

7.2 Second, it distinguishes local maturity from widened-profile immaturity

A text can be mature in its own domain. The graph may fall later because the profile widened scope, not because the original text never deserved peak.

7.3 Third, it keeps the extractor from over-owning the final judgment

The extractor does not decide peak. It only provides the ingredients. The profiler decides whether the ingredients are enough.

7.4 Fourth, it preserves inspectability

When a profile moves, the user can inspect why. Was it due to a gate closure, a knowledge failure, unresolved scope expansion, contradiction, or weak integration? This matters for the project's goal of truth-seeking rather than mystified authority.

8 Current Validation Program

The system is not yet a validated instrument. It is a serious candidate research program. The validation path should remain modular and staged.

8.1 Content validity

The first task is to establish whether the ontology and extraction schema actually capture the intended dimensions, gates, and profile-update signals. This requires expert review and careful challenge cases.

8.2 Extractor benchmarking

The extractor should be evaluated against adjudicated benchmark sets. The core question is not whether an LLM sounds philosophical, but whether it extracts the right structure under the fixed contract.

8.3 Inter-rater and model agreement

The system should compare human adjudicators, multiple extractor models, and repeated model runs. The aim is to measure whether the contract induces stable extraction rather than brittle artifacts.

8.4 Profiler stress testing

The deterministic layer should be stressed under weight perturbations, rule perturbations, and adversarial input forms. The system should show whether qualitative placement remains stable under reasonable implementation variance.

8.5 Predictive and longitudinal coherence

The final phases should ask whether diachronic profile movement is coherent over time. A profile should not drift wildly on semantically equivalent material, but it should visibly respond when new scope, contradiction, falsehood, or gate evidence is introduced.

9 Current Limitations

The system still has real limits.

First, extraction remains probabilistic. Even with a strong contract, bookkeeping mistakes and interpretation drift can happen. This is why some mechanical consistency checks belong in normalization code rather than prose alone.

Second, the present ontology is still being refined. It is already usable, but not yet frozen as a public instrument specification.

Third, gate relevance and scope behavior need sustained benchmark coverage. A system can appear stable on obvious tests while failing on rare but important adversarial cases.

Fourth, the system presently models text-conditioned worldview structure. It does not directly model the whole person independent of the available text and accumulated profile state.

10 Scientific Potential

If validated, Philosopher's Stone could become useful in several ways.

It could serve as a public reasoning aid, allowing people to test whether their new additions strengthen or weaken a profile. It could support dialogue by making disagreements more inspectable. It could give researchers a way to compare extractors, benchmark epistemic profiling systems, and study diachronic worldview development. It could also serve as a proof-of-concept that philosophical maturity can be modeled with explicit theory, constrained extraction, deterministic aggregation, and transparent geometry rather than sentiment buckets or opaque latent labels.

The broader value of the project is not that it claims final authority over human maturity. The value is that it tries to make philosophical evaluation auditable, challengeable, and scientifically tractable.

11 Conclusion

Philosopher's Stone is an attempt to build public infrastructure for modeling philosophical maturity and epistemic stability without hiding the logic of judgment. Its core commitments are explicit theory, structured extraction, deterministic profiling, and inspectable geometry.

The present version of the system is stronger than its earlier form because it now treats peak maturity as a scope-relative completion state rather than an artificially withheld asymptote. A text can reach objective peak when it genuinely contains the required ingredients. A profile can later fall when it widens scope and introduces error or incompleteness. That diachronic logic is not an implementation detail. It is central to what the system is trying to model.

The remaining work is empirical: freezing the ontology, benchmarking extractors, stress-testing the profiler, and validating the instrument. But the system is now much closer to a coherent research program than to a vague intuition. That is the point of the project.